Double Network Address Translation (NAT)

Sean McGrath
UCSC CMPE151 Spring 2015

What is Double NAT?

- When a router running network address translation is itself behind another NAT router.
- Normally implemented unknowingly in home or small business
- Example: connecting a wireless router to a modem provided by your ISP

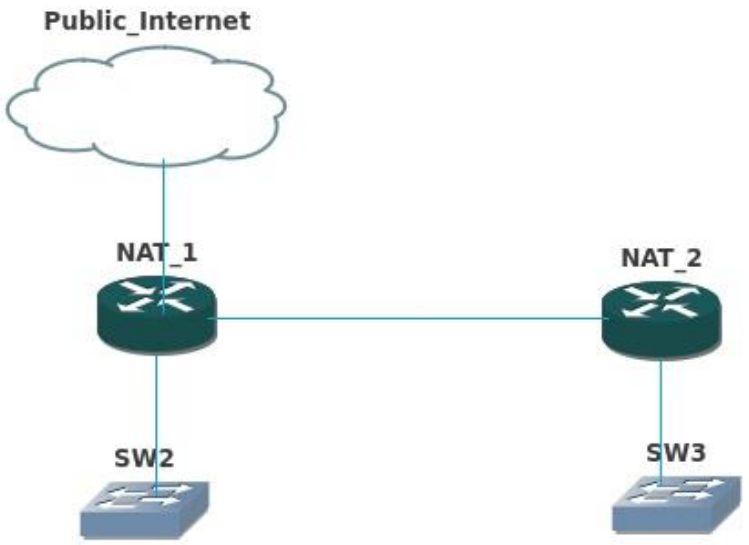

Goals of the Lab

- Successfully configure a double NAT topology
- Understand the address translation process while communicating between each subnet
- Determine if topology limits performance
- Discover how addressing problems can arise within the private network(s).

Topology 1

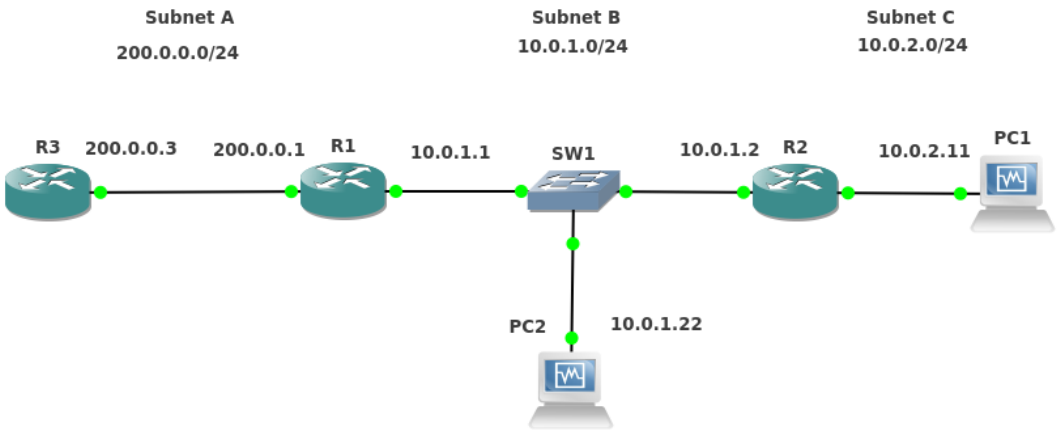

Pinging R3 from PC2

2	0.431554	10.0.1.22	200.0.0.3	ICMP	98	Echo (ping) request	id=0x3c17, seq=1/256, ttl=64
3	0.454504	200.0.0.3	10.0.1.22	ICMP	98	Echo (ping) reply	id=0x3c17, seq=1/256, ttl=254
4	1.436752	10.0.1.22	200.0.0.3	ICMP	98	Echo (ping) request	id=0x3c17, seq=2/512, ttl=64
5	1.465363	200.0.0.3	10.0.1.22	ICMP	98	Echo (ping) reply	id=0x3c17, seq=2/512, ttl=254

3	1.845607	200.0.0.1	200.0.0.3	ICMP	98	Echo (ping) request	id=0x3c17, seq=1/256, ttl=63
4	1.855670	200.0.0.3	200.0.0.1	ICMP	98	Echo (ping) reply	id=0x3c17, seq=1/256, ttl=255
5	2.856579	200.0.0.1	200.0.0.3	ICMP	98	Echo (ping) request	id=0x3c17, seq=2/512, ttl=63
6	2.866639	200.0.0.3	200.0.0.1	ICMP	98	Echo (ping) reply	id=0x3c17, seq=2/512, ttl=255

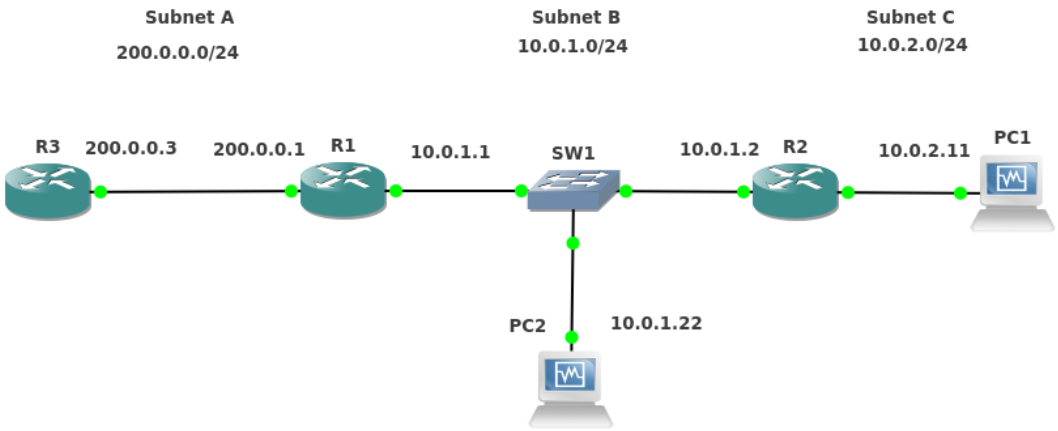

Issue a *ping* from PC1 to R3

5	20.545926	10.0.2.11	200.0.0.3	ICMP	98	Echo (ping) request id=0xda16, seq=1/256, ttl=64
6	20.592278	200.0.0.3	10.0.2.11	ICMP	98	Echo (ping) reply id=0xda16, seq=1/256, ttl=253
7	21.561705	10.0.2.11	200.0.0.3	ICMP	98	Echo (ping) request id=0xda16, seq=2/512, ttl=64
8	21.603735	200.0.0.3	10.0.2.11	ICMP	98	Echo (ping) reply id=0xda16, seq=2/512, ttl=253

3	12.451662	10.0.1.2	200.0.0.3	ICMP	98	Echo (ping) request id=0xda16, seq=1/256, ttl=63
4	12.481943	200.0.0.3	10.0.1.2	ICMP	98	Echo (ping) reply id=0xda16, seq=1/256, ttl=254
5	13.462919	10.0.1.2	200.0.0.3	ICMP	98	Echo (ping) request id=0xda16, seq=2/512, ttl=63
6	13.493386	200.0.0.3	10.0.1.2	ICMP	98	Echo (ping) reply id=0xda16, seq=2/512, ttl=254

3	3.872189	200.0.0.1	200.0.0.3	ICMP	98	Echo (ping) request id=0xda16, seq=1/256, ttl=62
4	3.882264	200.0.0.3	200.0.0.1	ICMP	98	Echo (ping) reply id=0xda16, seq=1/256, ttl=255
5	4.883511	200.0.0.1	200.0.0.3	ICMP	98	Echo (ping) request id=0xda16, seq=2/512, ttl=62
6	4.893582	200.0.0.3	200.0.0.1	ICMP	98	Echo (ping) reply id=0xda16, seq=2/512, ttl=255

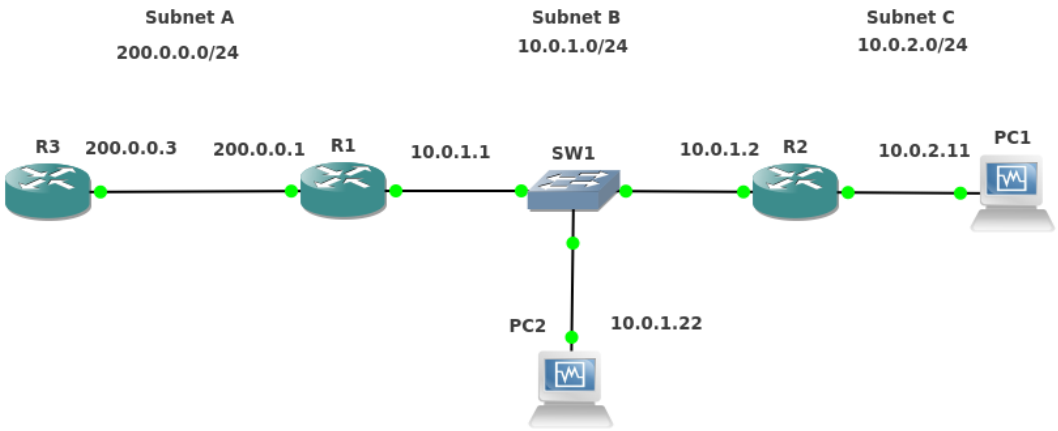

Subnet B and C are able to communicate since their IPv4 prefixes are necessarily different—R2 cannot have two interfaces with same IP prefix.

Examining Added Delay by 2nd NAT Router

- Simultaneously *pinged* R3 from PC1 and PC2

- Took average delay in milliseconds over 20 ICMP messages for each host
- Average delay for PC1 was about 20.25 ms more than PC2
- Suggests double NAT topology will slightly impact performance of streaming applications

	PC1	PC2
	65.6	69.9
	69.3	29.9
	49.9	31.9
	47.9	24.4
	43.2	56.3
	51.7	20.7
	67.6	31.4
	46.8	27.7
	48.4	23.2
	46.9	27
	43.1	30.8
	55.8	29.2
	45.6	28.1
	45.1	31
	84.3	36.8
	44.8	30.2
	42.1	54
	79.6	30.3
	43.8	27.9
	48.5	24.4
AVG	53.5	33.255
PC1-PC2	20.245	

Topology 2

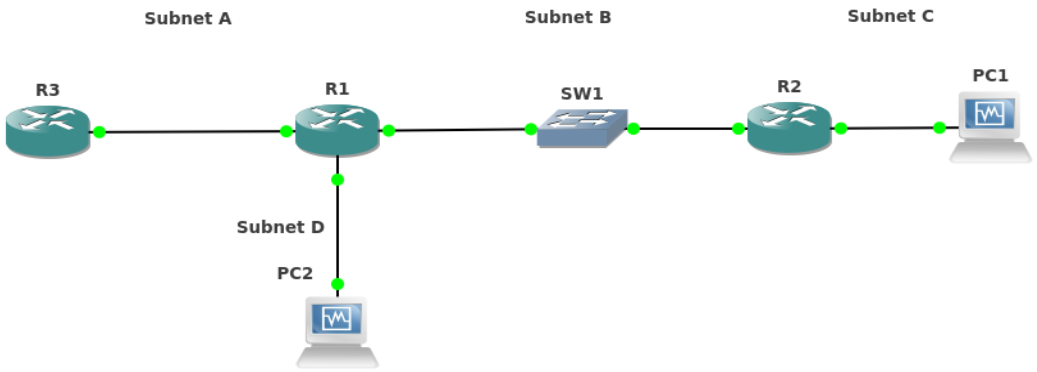

- PC2 is moved to a separate interface on R1, which is also NATed
- Configured Subnet D to have the *same* IP prefix as Subnet C

	Subnet A	Subnet B	Subnet C	Subnet D
IP prefix	200.0.0.0/24	10.0.1.0/24	10.0.2.0/24	10.0.2.0/24
R3	200.0.0.3			
R1	200.0.0.1	10.0.1.1		10.0.2.1
R2		10.0.1.2	10.0.2.2	
PC1			10.0.2.11	
PC2				10.0.2.22

Cannot communicate between Subnets C and D!

```
[root@PC2 ~]# ping -c3 10.0.2.11
PING 10.0.2.11 (10.0.2.11) 56(84) bytes of data.
From 10.0.2.22 icmp_seq=1 Destination Host Unreachable
From 10.0.2.22 icmp_seq=2 Destination Host Unreachable
From 10.0.2.22 icmp_seq=3 Destination Host Unreachable

--- 10.0.2.11 ping statistics ---
3 packets transmitted, 0 received, +3 errors, 100% packet loss, time 2009ms
, pipe 3
[root@PC2 ~]#
```

3	17.661092	CadmusCo_c4:88:b0	Broadcast	ARP	42	Who has 10.0.2.11? Tell 10.0.2.22
4	18.660738	CadmusCo_c4:88:b0	Broadcast	ARP	42	Who has 10.0.2.11? Tell 10.0.2.22
5	19.665678	CadmusCo_c4:88:b0	Broadcast	ARP	42	Who has 10.0.2.11? Tell 10.0.2.22

Solution: Static route to each host

```
C    200.0.0.0/24 is directly connected, FastEthernet0/0
     10.0.0.0/8 is variably subnetted, 3 subnets, 2 masks
S       10.0.2.11/32 [1/0] via 10.0.1.2
C       10.0.2.0/24 is directly connected, FastEthernet2/0
C       10.0.1.0/24 is directly connected, FastEthernet1/0
R1#
```

```
S    200.0.0.0/24 [1/0] via 10.0.1.1
     10.0.0.0/8 is variably subnetted, 3 subnets, 2 masks
C       10.0.2.0/24 is directly connected, FastEthernet1/0
C       10.0.1.0/24 is directly connected, FastEthernet0/0
S       10.0.2.22/32 [1/0] via 10.0.1.1
R2#show ip nat translations

R2#
```

Connectivity is restored!

Packet capture on R2's Subnet B interface

3	8.258486	10.0.2.22	10.0.2.11	ICMP	98	Echo (ping) request id=0x3116, seq=1/256, ttl=63
4	8.292208	10.0.1.2	10.0.2.22	ICMP	98	Echo (ping) reply id=0x3116, seq=1/256, ttl=63
5	9.240493	10.0.2.22	10.0.2.11	ICMP	98	Echo (ping) request id=0x3116, seq=2/512, ttl=63
6	9.260696	10.0.1.2	10.0.2.22	ICMP	98	Echo (ping) reply id=0x3116, seq=2/512, ttl=63

Packet capture on R2's Subnet C interface

2	8.167590	10.0.2.22	10.0.2.11	ICMP	98	Echo (ping) request id=0x3116, seq=1/256, ttl=62
5	8.179358	10.0.2.11	10.0.2.22	ICMP	98	Echo (ping) reply id=0x3116, seq=1/256, ttl=64
6	9.149647	10.0.2.22	10.0.2.11	ICMP	98	Echo (ping) request id=0x3116, seq=2/512, ttl=62
7	9.151762	10.0.2.11	10.0.2.22	ICMP	98	Echo (ping) reply id=0x3116, seq=2/512, ttl=64

	Subnet A	Subnet B	Subnet C	Subnet D
IP prefix	200.0.0.0/24	10.0.1.0/24	10.0.2.0/24	10.0.2.0/24
R3	200.0.0.3			
R1	200.0.0.1	10.0.1.1		10.0.2.1
R2		10.0.1.2	10.0.2.2	
PC1			10.0.2.11	
PC2				10.0.2.22

A Subtle Potential for Problems…

- Add a second host onto Subnet D and give it the same IP address as the host on Subnet C
- Now PC2 cannot communicate directly with PC1, PC2 sends the message to PC3.
- However, communication can happen if PC1 initiates the message to PC2!

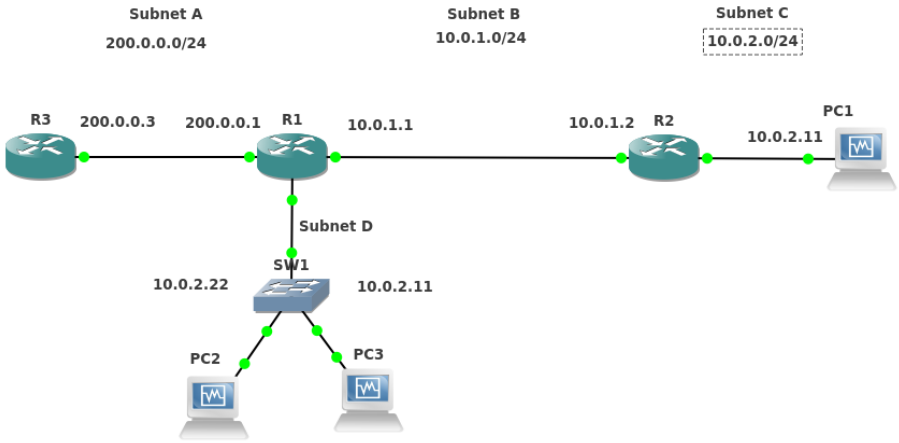

1	0.000000	cc:00:6f:97:00:20	cc:00:6f:97:00:20	LOOP	60	Reply
2	5.447980	10.0.1.2	10.0.2.22	ICMP	98	Echo (ping) request id=0x2517, seq=1/256, ttl=62
3	5.451736	CadmusCo_c4:88:b0	Broadcast	ARP	42	Who has 10.0.2.1? Tell 10.0.2.22
4	5.518657	cc:00:6f:97:00:20	CadmusCo_c4:88:b0	ARP	60	10.0.2.1 is at cc:00:6f:97:00:20
5	5.519202	10.0.2.22	10.0.1.2	ICMP	98	Echo (ping) reply id=0x2517, seq=1/256, ttl=64
6	6.417684	10.0.1.2	10.0.2.22	ICMP	98	Echo (ping) request id=0x2517, seq=2/512, ttl=62
7	6.418128	10.0.2.22	10.0.1.2	ICMP	98	Echo (ping) reply id=0x2517, seq=2/512, ttl=64

Problems Encountered

- NAT overload (not in netref)
- This is equivalent to Port Address Translation—allows multiple internal addresses to be mapped to a single public address
- Simple solution

ip access-list standard NAT
permit 10.0.1.0 0.0.0.255
ip nat inside source list NAT interface fastEthernet 0/1 overload

Conclusions

- Double NAT isn't "evil", if it is configured well then it works, because NAT works!
- Performance impact is minimal
- Problem with topology is that people are usually unaware that it's been implemented—makes troubleshooting a challenge
- Can be a simple way of segregating different parts of your network

NAT overload source

http://www.cisco.com/c/en/us/support/docs/ip/network-address-translation-nat/13772-12.html

Questions?...